

Spark assignment

- Author: David Kozak
- Email: dkozak94@gmail.com

The goal of this assignment was to implement 12 data analysis tasks using [Apache Spark](#). The structure of the report is as follows. First, the dataset is introduced. Afterwards descriptions and solutions for each of the tasks are given.

Dataset

As in the first [mongo assignment](#), I decided to use the dataset [employee reviews](#), which contains reviews from employees of 6 IT companies - Google, Facebook, Amazon, Apple and Netflix. The reviews are stored in a [csv file](#). The structure of one review can be determined by checking the header line.

```
head -n 1 employee_reviews.csv
```

Which results in:

```
id,company,location,dates,job-title,summary,pros,cons,advice-to-mgmt,overall-ratings,work-balance-stars,culture-values-stars,carrer-opportunities-stars,comp-benefit-stars,senior-mangemnet-stars,helpful-count,link
```

The size is 67529 reviews.

```
tail -n 1
```

```
67529,microsoft,none," Dec 14, 2010",Former Employee - Sen ....
```

Run the code

The whole repository is a gradle module and the source code is written in [Kotlin](#). It can be compiled into jar using the following command.

```
gradle jar
```

This jar can then be submitted to locally installed apache spark instance.

```
${SPARK_EXEC} --class io.dkozak.estg.spark.assignment.MainKt \
--master "local[4]" ${JAR_FILE} ${INPUT_FILE} ${OUTPUT_DIR}
```

- SPARK_EXEC is the location of the spark-submit file in the bin folder inside the spark directory
- JAR_FILE is the jar build the previous step
- INPUT_FILE is a path to the input csv file
- OUTPUT_DIR is a directory into which the output should be written

To allow for easier execution, [execute.sh](#) shell script is prepared. This script takes following arguments.

- location where ApacheSpark is installed
- [optional] a single task that should be run (default is all tasks)

Please note that for this script to work correctly, the input csv file has to be stored at path `./employee_reviews.csv`.

The output can be found in the `./output` directory, each task results are stored in a subfolder identified by the number of the task.

Tasks

In this section description and solutions for individual tasks will be given

1) [Lookup collection](#)

The goal of this task was to replace one enumeration column in the dataset with a numeric value. The mapping between the numeric value and the original enumeration value should be saved in a separate lookup collection.

For this task I decided to replace the company column in the original dataset.

First I created the lookup collection, saved it and loaded it into map.

```
private fun prepareCompanyDataset(
    dataset: Dataset<Row>,
    outputDir: String
): Dataset<Row> {
    val companyDataset = dataset.select("company")
        .distinct()
        .coalesce(1) dataset.col("company")
    var whenColumn: Column? = null
    for ((company, id) in companies) {
        whenColumn = if (whenColumn == null) {
            `when` (companyColumn.equalTo(company), id)
        } else {
            whenColumn.`when` (companyColumn.equalTo(company), id)
        }
    }
    companyDataset
        .withColumn("id", monotonically_increasing_id())
        .writeCsv("$outputDir/lookup")
    return companyDataset
}

val companies = prepareCompanyDataset(dataset, outputDir)
    .collectAsList()
    .map { it.getString(0) }
    .mapIndexed { i, company ->
        company to i
    }.toMap()
```

Then I used this map to generate a query that replaces the column values in the original dataset with numerical values.

```
val companyColumn = dataset.col("company")
var whenColumn: Column? = null
for ((company, id) in companies) {
    whenColumn = if (whenColumn == null) {
        `when` (companyColumn.equalTo(company), id)
    } else {
        whenColumn.`when` (companyColumn.equalTo(company), id)
    }
}

dataset.withColumn("company", whenColumn)
    .writeCsv("$outputDir/reviews")
```

2) Oversampling

Since the amount of reviews per company is different, it might be important for some types of analysis to transform the dataset in a way to make sure that all the categories are represented equally. Therefore, for the oversampling task, I decided to even out the amount of reviews per company. I achieved this by inserting some reviews multiple times.

First I calculated how many reviews per company are in the dataset.

```
val reviewsPerCompany = dataset.groupBy("company")
    .count()
    .orderBy(desc("count"))
    .collectAsList()
    .map { it.getString(0) to it.getLong(1) }

val (maxCompanyName, max) = reviewsPerCompany[0]
```

Then I used this to compute how many reviews need to be added.

```
val reviewsToAdd = reviewsPerCompany
    .filter { (name, _) -> name != maxCompanyName }
    .map { (name, count) -> Triple(name, count, max - count) }
```

In the end I used this information to generate a new bigger dataset in which all companies are represented equally.

```

var result = dataset.where(dataset.col("company").equalTo(maxCompanyName))
for ((name, count, toAdd) in reviewsToAdd) {
  logger.task("Oversampling on $name") {
    val reviewsForCompany = dataset.where(dataset.col("company").equalTo(name))

    var buffer = reviewsForCompany
    val fullIterations = toAdd / count
    val reminder = toAdd % count
    val percentage = reminder.toDouble() / count
    logger.log("Need $fullIterations full iterations and another $percentage in the last one")
    for (i in 0 until fullIterations) {
      buffer = buffer.union(reviewsForCompany)
    }
    if (reminder > 0) {
      val subset = reviewsForCompany.sample(percentage)
      buffer = buffer.union(subset)
    }

    logger.log("Final review count for $name is ${buffer.count()}")
    result = result.union(buffer)
  }
}

```

3) Undersampling

As in the case of mongodb, undersampling was easier to implement, because there is a sample operator that takes a randomly selected subset from the dataset. Again, I used undersampling to ensure that the amount of tasks per company is equal.

The initial procedures were almost the same as the previous task, except that I wanted to find out the company with lowest amount of reviews, so that I could sample the reviews from other companies accordingly.

```

var result = dataset.where(dataset.col("company").equalTo(minCompanyName))
for ((companyName, percentage) in percentagePerCompany) {
  val sampled = dataset.where(dataset.col("company").equalTo(companyName))
    .sample(percentage)
  result = result.union(sampled)
}

```

4) Discretizing

For this task I decided to discretize the overall rating column. Originally it contained values from the interval <0.0,5.0> and I transformed it into values from {1,2,3,4,5} by rounding the floating point numbers up to their nearest integer.

```

val col = dataset.col("overall-ratings")
val columnRule = `when`(col.leq(1), 1)
  .`when`(col.leq(2), 2)
  .`when`(col.leq(3), 3)
  .`when`(col.leq(4), 4)
  .otherwise(5)

val discretized = dataset.withColumn("overall-ratings", columnRule)

```

5) Probabilistic analysis

For this task I decided to calculate the probabilities of a review belonging to a company. For that I counted how many reviews are in the whole dataset and then for each company I divided the amount of it's review by the amount of all reviews.

```

val totalReviews = dataset.count().toDouble()
val probabilities = dataset.groupBy("company")
  .count()
  .collectAsList()
  .map { it.getString(0) to it.getLong(1) }
  .map { (name, count) -> name to count / totalReviews }

```

6) Tf-idf

As in the mongodb task, I decided to calculate the tf-idf for the word 'work' in the summary column.

First, I had to count how many documents contain that word.

```
val containsWord = object : Function1<Row, Int>, Serializable {
    override fun apply(it: Row): Int {
        return if ((it.getString(summaryIndex) ?: "").contains(" $word ")) 1 else 0
    }
}

val count = dataset.map(containsWord, Encoders.INT())
    .agg(sum("value")).collectAsList()[0].getLong(0).toDouble()
```

From that, I could compute idf.

```
val idf = Math.log(dataset.count() / count)
```

Afterwards, I computed tf-idf for each document.

```
val idfCalc = object : Function1<Row, Double>, Serializable {
    override fun apply(it: Row): Double {
        val words = (it.getString(summaryIndex) ?: "").split(" ")
        val workCount = words.count { it == word }.toDouble()
        val tf = workCount / words.size
        return tf * idf
    }
}

val tf_idf = dataset.map(idfCalc, Encoders.DOUBLE())
```

7) Index

Since the column containing longest strings is the summary, I decided to create inverted index for this one.

First, I transformed documents into tuples, where the first element is a word and the second element is a documentId of a document, where this word can be found.

```
val getText = object : FlatMapFunction<Row, Tuple2<String, String>>, Serializable {

    override fun call(it: Row): MutableIterator<Tuple2<String, String>> {
        val id = it.getString(idIndex)
        val text = (it.getString(summaryIndex) ?: "").
            .split(" ")
            .filter { it.isNotEmpty() }
            .map { Tuple2(it, id) }
        return object : MutableIterator<Tuple2<String, String>> {

            var i = 0

            override fun hasNext(): Boolean = i < text.size

            override fun next(): Tuple2<String, String> = text[i++]
        }
    }
}
```

Then I grouped them by the first column, that is the word, merging all the document ids together.

```
val result = dataset.flatMap(getText, Encoders.tuple(Encoders.STRING(), Encoders.STRING()))
    .groupBy("_1")
    .agg(functions.concat_ws(" ", collect_list("_2")))
    .toDF("word", "documents")
```

8) k-fold cross validation

For this task I decided to split the dataset into 5 disjoint datasets that can be used for example to test that a prediction model is not overfitted.

For that I iterated over the original dataset and sliced it.

```
val foldSize = (dataset.count() / k).toInt()

for (i in 0 until k) {
    val start = i * foldSize
```

```

val rows = dataset
  .where(dataset.col("id").`$greater$eq`(start))
  .limit(foldSize)

rows.show()
rows.writeCsv("$outputDir/fold_$i")
}

```

9) Normalization

I decided to normalize the values in overall-ratings, which originally were form the interval <0.0,5.0>.

First I calculated new values in a temporal dataset.

```

val normalizationFunction = object : Function1<Row, Double>, Serializable {
  override fun apply(it: Row): Double {
    val rating = it.getString(colIndex).toDoubleOrNull()
    ?: throw IllegalArgumentException("Could not parse ${it.getString(colIndex)}")
    return rating / 5
  }
}

val resultCol = dataset.map(normalizationFunction, Encoders.DOUBLE())
  .withColumn("id", functions.monotonically_increasing_id())

```

Then I merged it back to the original one.

```

val joined = dataset
  .drop("overall-ratings")
  .join(resultCol, "id")

```

10) Remove noise

I decided to remove all reviews that were older than 1.1.2017, because they are outdated and therefore they are less valuable for people deciding in which company to go now.

```

val filtered = dataset.filter {
  val dateString = it.getString(dateColIndex).trim()
  try {
    val dateFormat = DateTimeFormatter.ofPattern("MMM d, yyyy")
    val date = LocalDate.parse(dateString, dateFormat)
    date >= limit
  } catch (ex: DateTimeParseException) {
    false
  }
}

```

11) Fill missing values

Again I encountered the problem that I could not find missing values. However, to fulfil the task, I decided to create a code that would insert average rating into reviews where overall-rating would be missing.

```

val avgRating = dataset.agg(avg("overall-ratings"))
  .collectAsList()[0]
  .getDouble(0)

val addAverage = `when`(dataset.col("overall-ratings").isNull, avgRating)

dataset.withColumn("overall-ratings", addAverage)

```

12) Pivot table

For this tasks, I decided to compute the average value per company for the following columns. * overall-ratings * work-balance-stars * culture-values-stars * carrer-opportunities-stars

```

val pivot = dataset
  .groupBy("company")
  .agg(

```

```
avg("overall-ratings"),  
avg("work-balance-stars"),  
avg("culture-values-stars"),  
avg("carrer-opportunities-stars")  
)
```